

Analysis of unstructured data

Lecture 4 - pandas examples

Janusz Szwabiński

Overview:

- Reading from CSV files
- Groupby and aggregate
- Searching for information
- Data cleaning
- String operations
- Parsing Unix timestamps
- Loading data from SQL
- Pandas vs SQL

Referencje:

- Homepage of pandas project: <http://pandas.pydata.org/> (<http://pandas.pydata.org/>)
- Pandas Cookbook, <https://github.com/jvns/pandas-cookbook> (<https://github.com/jvns/pandas-cookbook>)

Data sources:

- <http://donnees.ville.montreal.qc.ca/dataset/velos-comptage> (<http://donnees.ville.montreal.qc.ca/dataset/velos-comptage>) - cyclist data from Montreal
- http://climate.weather.gc.ca/index_e.html (http://climate.weather.gc.ca/index_e.html) - Canadian weather data
- <https://nycopendata.socrata.com/> (<https://nycopendata.socrata.com/>) - 311 service requests from NYC

In [1]:

```
%matplotlib inline
import numpy as np
import pandas as pd
```

Reading from CSV files

In [2]:

```
! head data/bikes.csv
```

In [3]:

```
broken_df = pd.read_csv('data/bikes.csv')  
broken_df[:3]
```

```

-----
UnicodeDecodeError                                Traceback (most recent call last)
<ipython-input-3-823d15dd87ce> in <module>()
----> 1 broken_df = pd.read_csv('data/bikes.csv')
      2 broken_df[:3]

/usr/local/lib/python3.5/dist-packages/pandas/io/parsers.py in parse
r_f(filepath_or_buffer, sep, delimiter, header, names, index_col, usecols, squeeze, prefix, mangle_dupe_cols, dtype, engine, converters, true_values, false_values, skipinitialspace, skiprows, nrows, na_values, keep_default_na, na_filter, verbose, skip_blank_lines, parse_dates, infer_datetime_format, keep_date_col, date_parser, dayfirst, iterator, chunksize, compression, thousands, decimal, lineterminator, quotechar, quoting, escapechar, comment, encoding, dialect, tupleize_cols, error_bad_lines, warn_bad_lines, skipfooter, skip_footer, doquote, delim_whitespace, as_recarray, compact_ints, use_unsigned, low_memory, buffer_lines, memory_map, float_precision)
    643         skip_blank_lines=skip_blank_lines)
    644
--> 645         return _read(filepath_or_buffer, kwds)
    646
    647     parser_f.__name__ = name

/usr/local/lib/python3.5/dist-packages/pandas/io/parsers.py in
_read(filepath_or_buffer, kwds)
    386
    387     # Create the parser.
--> 388     parser = TextFileReader(filepath_or_buffer, **kwds)
    389
    390     if (nrows is not None) and (chunksize is not None):

/usr/local/lib/python3.5/dist-packages/pandas/io/parsers.py in __init__(self, f, engine, **kwds)
    727         self.options['has_index_names'] = kwds['has_index_names']
    728
--> 729         self._make_engine(self.engine)
    730
    731     def close(self):

/usr/local/lib/python3.5/dist-packages/pandas/io/parsers.py in _make_engine(self, engine)
    920     def _make_engine(self, engine='c'):
    921         if engine == 'c':
--> 922             self._engine = CParserWrapper(self.f, **self.options)
    923         else:
    924             if engine == 'python':

/usr/local/lib/python3.5/dist-packages/pandas/io/parsers.py in __init__(self, src, **kwds)
   1387         kwds['allow_leading_cols'] = self.index_col is not False
   1388
-> 1389         self._reader = _parser.TextReader(src, **kwds)
   1390
   1391         # XXX

```

pandas/parser.pyx in pandas.parser.TextReader.__cinit__ (pandas/pars

```
er.c:6077)()
```

```
pandas/parser.pyx in pandas.parser.TextReader._get_header (pandas/pa
rser.c:9215)()
```

```
UnicodeDecodeError: 'utf-8' codec can't decode byte 0xe9 in position
15: invalid continuation byte
```

```
In [4]:
```

```
# we set the separator to ';' and change the encoding to latin1
# we want to parse the dates
# we want the dates to have the day first
# we set the index to be the date column
fixed_df = pd.read_csv('data/bikes.csv', sep=';', encoding='latin1',
parse_dates=['Date'], dayfirst=True, index_col='Date')
fixed_df[:3]
```

```
Out[4]:
```

	Berri 1	Brébeuf (données non disponibles)	Côte- Sainte- Catherine	Maisonneuve 1	Maisonneuve 2	du Parc	Pierre- Dupuy	Rach
Date								
2012- 01-01	35	NaN	0	38	51	26	10	16
2012- 01-02	83	NaN	1	68	153	53	6	43
2012- 01-03	135	NaN	2	104	248	89	3	58

Selecting a column

In [5]:

```
fixed_df['Berri 1']
```

Out[5]:

Date	
2012-01-01	35
2012-01-02	83
2012-01-03	135
2012-01-04	144
2012-01-05	197
2012-01-06	146
2012-01-07	98
2012-01-08	95
2012-01-09	244
2012-01-10	397
2012-01-11	273
2012-01-12	157
2012-01-13	75
2012-01-14	32
2012-01-15	54
2012-01-16	168
2012-01-17	155
2012-01-18	139
2012-01-19	191
2012-01-20	161
2012-01-21	53
2012-01-22	71
2012-01-23	210
2012-01-24	299
2012-01-25	334
2012-01-26	306
2012-01-27	91
2012-01-28	80
2012-01-29	87
2012-01-30	219
...	
2012-10-07	1580
2012-10-08	1854
2012-10-09	4787
2012-10-10	3115
2012-10-11	3746
2012-10-12	3169
2012-10-13	1783
2012-10-14	587
2012-10-15	3292
2012-10-16	3739
2012-10-17	4098
2012-10-18	4671
2012-10-19	1313
2012-10-20	2011
2012-10-21	1277
2012-10-22	3650
2012-10-23	4177
2012-10-24	3744
2012-10-25	3735
2012-10-26	4290
2012-10-27	1857
2012-10-28	1310
2012-10-29	2919
2012-10-30	2887
2012-10-31	2634
2012-11-01	2405
2012-11-02	1582
2012-11-03	811

```
2012-11-03    877  
2012-11-04    966  
2012-11-05   2247  
Name: Berri 1, dtype: int64
```


Plotting a column

In [6]:

```
fixed_df['Berri 1'].plot()
```

Out[6]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19c341bda0>




Plotting all columns

In [7]:

```
fixed_df.plot(figsize=(15, 10))
```

Out[7]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f19c337bda0>
```



Groupby and aggregate

In [8]:

```
berri_bikes = fixed_df[['Berri 1']].copy()
berri_bikes[:5]
```

Out[8]:

	Berri 1
Date	
2012-01-01	35
2012-01-02	83
2012-01-03	135
2012-01-04	144
2012-01-05	197

In [9]:

```
berri_bikes.index
```

Out[9]:

```
DatetimeIndex(['2012-01-01', '2012-01-02', '2012-01-03', '2012-01-04',  
              '2012-01-05', '2012-01-06', '2012-01-07', '2012-01-08',  
              '2012-01-09', '2012-01-10',  
              ...,  
              '2012-10-27', '2012-10-28', '2012-10-29', '2012-10-30',  
              '2012-10-31', '2012-11-01', '2012-11-02', '2012-11-03',  
              '2012-11-04', '2012-11-05'],  
              dtype='datetime64[ns]', name='Date', length=310, freq=None)
```

In [10]:

berri_bikes.index.day

Out[10]:

[illegible]

Adding a weekday column to the dataframe

In [11]:

```
berri_bikes.index.weekday # 0 - Monday
```

Out[11]:

```
array([6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4,
       5, 6, 0,
       1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6,
       0, 1, 2,
       3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1,
       2, 3, 4,
       5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3,
       4, 5, 6,
       0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5,
       6, 0, 1,
       2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0,
       1, 2, 3,
       4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2,
       3, 4, 5,
       6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4,
       5, 6, 0,
       1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6,
       0, 1, 2,
       3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1,
       2, 3, 4,
       5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3,
       4, 5, 6,
       0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5,
       6, 0, 1,
       2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0,
       1, 2, 3,
       4, 5, 6, 0, 1, 2, 3, 4, 5, 6, 0], dtype=int32)
```

In [12]:

```
berri_bikes.loc[:, 'weekday'] = berri_bikes.index.weekday
berri_bikes[:5]
```

Out[12]:

	Berri 1	weekday
Date		
2012-01-01	35	6
2012-01-02	83	0
2012-01-03	135	1
2012-01-04	144	2
2012-01-05	197	3

Adding up the cyclists by weekday

In [13]:

```
weekday_counts = berri_bikes.groupby('weekday').aggregate(sum)
weekday_counts
```

Out[13]:

	Berri 1
weekday	
0	134298
1	135305
2	152972
3	160131
4	141771
5	101578
6	99310

In [14]:

```
weekday_counts.index = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',  
                        'Saturday', 'Sunday']  
weekday_counts
```

Out[14]:


	Berri 1
Monday	134298
Tuesday	135305
Wednesday	152972
Thursday	160131
Friday	141771
Saturday	101578
Sunday	99310

In [15]:

```
weekday_counts.plot(kind='bar')
```

Out[15]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19c3385cc0>



It looks like Montrealers are commuter cyclists - they bike much more during the week!

Searching for information

We're going to use a larger dataset - a subset of 311 service requests from New York City:

In [16]:

```
complaints = pd.read_csv('data/311-service-requests.csv')
```

```
/usr/local/lib/python3.5/dist-packages/IPython/core/interactiveshell  
l.py:2698: DtypeWarning: Columns (8) have mixed types. Specify dtype  
option on import or set low_memory=False.  
  interactivity=interactivity, compiler=compiler, result=result)
```

In [17]:

```
complaints = pd.read_csv('data/311-service-requests.csv', low_memory=False)
```

In [18]:

```
complaints
```

Out[18]:

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	De
0	26589651	10/31/2013 02:08:41 AM	NaN	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot
1	26593698	10/31/2013 02:01:04 AM	NaN	NYPD	New York City Police Department	Illegal Parking	Co Ov
2	26594139	10/31/2013 02:00:24 AM	10/31/2013 02:40:32 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
3	26595721	10/31/2013 01:56:23 AM	10/31/2013 02:21:48 AM	NYPD	New York City Police Department	Noise - Vehicle	Ca
4	26590930	10/31/2013 01:53:44 AM	NaN	DOHMH	Department of Health and Mental Hygiene	Rodent	Co Att Ro
5	26592370	10/31/2013 01:46:52 AM	NaN	NYPD	New York City Police Department	Noise - Commercial	Ba
6	26595682	10/31/2013 01:46:40 AM	NaN	NYPD	New York City Police Department	Blocked Driveway	No
7	26595195	10/31/2013 01:44:19 AM	10/31/2013 01:58:49 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
8	26590540	10/31/2013 01:44:14 AM	10/31/2013 02:28:04 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
9	26594392	10/31/2013 01:34:41 AM	10/31/2013 02:23:51 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
10	26595176	10/31/2013 01:25:12 AM	NaN	NYPD	New York City Police Department	Noise - House of Worship	Lot
11	26591982	10/31/2013 01:24:14 AM	10/31/2013 01:54:39 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
12	26594169	10/31/2013 01:20:57 AM	10/31/2013 02:12:31 AM	NYPD	New York City Police Department	Illegal Parking	Do Blo


	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	De
13	26594391	10/31/2013 01:20:13 AM	NaN	NYPD	New York City Police Department	Noise - Vehicle	En
14	26590917	10/31/2013 01:19:54 AM	NaN	DOHMH	Department of Health and Mental Hygiene	Rodent	Ra
15	26591458	10/31/2013 01:14:02 AM	10/31/2013 01:30:34 AM	NYPD	New York City Police Department	Noise - House of Worship	Lot
16	26594086	10/31/2013 12:54:03 AM	10/31/2013 02:16:39 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot
17	26595117	10/31/2013 12:52:46 AM	NaN	NYPD	New York City Police Department	Illegal Parking	Pos Sig
18	26590389	10/31/2013 12:51:00 AM	NaN	DOT	Department of Transportation	Street Light Condition	Str
19	26594210	10/31/2013 12:46:27 AM	NaN	NYPD	New York City Police Department	Noise - Commercial	Lot
20	26592932	10/31/2013 12:43:47 AM	10/31/2013 12:56:20 AM	NYPD	New York City Police Department	Noise - House of Worship	Lot
21	26594152	10/31/2013 12:41:17 AM	10/31/2013 01:04:37 AM	NYPD	New York City Police Department	Noise - Commercial	Bal
22	26589678	10/31/2013 12:39:55 AM	NaN	NYPD	New York City Police Department	Noise - Vehicle	Ca
23	26592304	10/31/2013 12:38:00 AM	NaN	NYPD	New York City Police Department	Noise - Commercial	Lot
24	26591892	10/31/2013 12:37:16 AM	NaN	NYPD	New York City Police Department	Blocked Driveway	Pa
25	26591573	10/31/2013 12:35:18 AM	10/31/2013 02:41:35 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	De
26	26590509	10/31/2013 12:33:00 AM	NaN	DOT	Department of Transportation	Street Light Condition	Str
27	26591379	10/31/2013 12:32:44 AM	NaN	DOHMH	Department of Health and Mental Hygiene	Harboring Bees/Wasps	Be a b
28	26594085	10/31/2013 12:32:08 AM	NaN	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot
29	26589201	10/31/2013 12:32:00 AM	NaN	DOT	Department of Transportation	Street Light Condition	Str
...
111039	26428764	10/04/2013 12:17:03 AM	10/04/2013 12:38:37 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
111040	26426166	10/04/2013 12:16:22 AM	10/04/2013 05:50:49 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
111041	26438565	10/04/2013 12:16:00 AM	NaN	DEP	Department of Environmental Protection	Noise	No Co Bel Ho
111042	26428990	10/04/2013 12:15:52 AM	10/04/2013 12:44:52 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot
111043	26432659	10/04/2013 12:15:46 AM	10/04/2013 04:18:45 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
111044	26426096	10/04/2013 12:14:09 AM	10/04/2013 01:03:46 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot
111045	26437764	10/04/2013 12:14:00 AM	10/04/2013 12:14:00 AM	DEP	Department of Environmental Protection	Water System	Dir
111046	26436286	10/04/2013 12:14:00 AM	NaN	DEP	Department of Environmental Protection	Noise	No Co Bel Ho

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	De
111047	26428989	10/04/2013 12:13:08 AM	10/04/2013 02:12:47 AM	NYPD	New York City Police Department	Illegal Parking	Pos Sig
111048	26430030	10/04/2013 12:12:07 AM	10/04/2013 02:45:24 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot
111049	26429663	10/04/2013 12:12:07 AM	10/04/2013 01:03:44 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
111050	26437763	10/04/2013 12:11:00 AM	NaN	DEP	Department of Environmental Protection	Noise	No Co Bel Ho
111051	26432955	10/04/2013 12:08:15 AM	10/04/2013 12:48:02 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
111052	26437035	10/04/2013 12:08:00 AM	10/04/2013 12:13:00 AM	DEP	Department of Environmental Protection	Water System	Dir
111053	26433197	10/04/2013 12:08:00 AM	10/04/2013 12:00:00 PM	DSNY	BCC - Queens East	Derelict Vehicles	14 Vel
111054	26426060	10/04/2013 12:06:39 AM	10/04/2013 12:31:16 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot
111055	26430628	10/04/2013 12:06:28 AM	10/04/2013 12:21:39 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
111056	26431648	10/04/2013 12:06:26 AM	10/23/2013 08:14:52 AM	DOT	Department of Transportation	Street Sign - Missing	Bu:
111057	26437034	10/04/2013 12:06:00 AM	NaN	DEP	Department of Environmental Protection	Noise	No Ha
111058	26426094	10/04/2013 12:05:12 AM	10/04/2013 01:08:29 AM	NYPD	New York City Police Department	Noise - Commercial	Lot
111059	26429040	10/04/2013 12:04:52 AM	10/04/2013 03:01:04 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lot

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	De
111060	26434084	10/04/2013 12:04:00 AM	NaN	DEP	Department of Environmental Protection	Noise	No Co Bel Ho
111061	26426164	10/04/2013 12:03:00 AM	10/04/2013 02:14:57 AM	NYPD	New York City Police Department	Noise - Commercial	Lo
111062	26439710	10/04/2013 12:03:00 AM	10/04/2013 12:03:00 AM	DEP	Department of Environmental Protection	Water System	Dir
111063	26435569	10/04/2013 12:02:00 AM	10/04/2013 01:10:00 AM	DEP	Department of Environmental Protection	Water System	Dir
111064	26426013	10/04/2013 12:01:13 AM	10/07/2013 04:07:16 PM	DPR	Department of Parks and Recreation	Maintenance or Facility	Str Ou
111065	26428083	10/04/2013 12:01:05 AM	10/04/2013 02:13:50 AM	NYPD	New York City Police Department	Illegal Parking	Pos Sig
111066	26428987	10/04/2013 12:00:45 AM	10/04/2013 01:25:01 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Lo
111067	26426115	10/04/2013 12:00:28 AM	10/04/2013 04:17:32 AM	NYPD	New York City Police Department	Noise - Commercial	Lo
111068	26428033	10/04/2013 12:00:10 AM	10/04/2013 01:20:52 AM	NYPD	New York City Police Department	Blocked Driveway	Pa

111069 rows × 52 columns



Selecting columns and rows

In [19]:

```
complaints['Complaint Type']
```

Out[19]:

```

0      Noise - Street/Sidewalk
1      Illegal Parking
2      Noise - Commercial
3      Noise - Vehicle
4      Rodent
5      Noise - Commercial
6      Blocked Driveway
7      Noise - Commercial
8      Noise - Commercial
9      Noise - Commercial
10     Noise - House of Worship
11     Noise - Commercial
12     Illegal Parking
13     Noise - Vehicle
14     Rodent
15     Noise - House of Worship
16     Noise - Street/Sidewalk
17     Illegal Parking
18     Street Light Condition
19     Noise - Commercial
20     Noise - House of Worship
21     Noise - Commercial
22     Noise - Vehicle
23     Noise - Commercial
24     Blocked Driveway
25     Noise - Street/Sidewalk
26     Street Light Condition
27     Harboring Bees/Wasps
28     Noise - Street/Sidewalk
29     Street Light Condition

...
111039    Noise - Commercial
111040    Noise - Commercial
111041    Noise
111042    Noise - Street/Sidewalk
111043    Noise - Commercial
111044    Noise - Street/Sidewalk
111045    Water System
111046    Noise
111047    Illegal Parking
111048    Noise - Street/Sidewalk
111049    Noise - Commercial
111050    Noise
111051    Noise - Commercial
111052    Water System
111053    Derelict Vehicles
111054    Noise - Street/Sidewalk
111055    Noise - Commercial
111056    Street Sign - Missing
111057    Noise
111058    Noise - Commercial
111059    Noise - Street/Sidewalk
111060    Noise
111061    Noise - Commercial
111062    Water System
111063    Water System
111064    Maintenance or Facility
111065    Illegal Parking
111066    Noise - Street/Sidewalk
111067    Noise - Commercial

```

111067

NOISE - COMMERCIAL

111068 Blocked Driveway
 Name: Complaint Type, dtype: object


In [20]:

complaints[:5]

Out[20]:

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	Descriptor
0	26589651	10/31/2013 02:08:41 AM	NaN	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Talking
1	26593698	10/31/2013 02:01:04 AM	NaN	NYPD	New York City Police Department	Illegal Parking	Commercial Overnight Parking
2	26594139	10/31/2013 02:00:24 AM	10/31/2013 02:40:32 AM	NYPD	New York City Police Department	Noise - Commercial	Loud Music/Party
3	26595721	10/31/2013 01:56:23 AM	10/31/2013 02:21:48 AM	NYPD	New York City Police Department	Noise - Vehicle	Car/Truck Horn
4	26590930	10/31/2013 01:53:44 AM	NaN	DOHMH	Department of Health and Mental Hygiene	Rodent	Condition Attracting Rodents

5 rows × 52 columns



In [21]:

complaints['Complaint Type'][:5]

Out[21]:

```
0    Noise - Street/Sidewalk
1         Illegal Parking
2         Noise - Commercial
3         Noise - Vehicle
4             Rodent
Name: Complaint Type, dtype: object
```

Selecting multiple columns

In [22]:

```
complaints[['Complaint Type', 'Borough']][:5]
```

Out[22]:

	Complaint Type	Borough
0	Noise - Street/Sidewalk	QUEENS
1	Illegal Parking	QUEENS
2	Noise - Commercial	MANHATTAN
3	Noise - Vehicle	MANHATTAN
4	Rodent	MANHATTAN

What is the most common complaint type?

In [23]:

```
complaints['Complaint Type'].value_counts()
```

Out[23]:

HEATING	14200
GENERAL CONSTRUCTION	7471
Street Light Condition	7117
DOF Literature Request	5797
PLUMBING	5373
PAINT - PLASTER	5149
Blocked Driveway	4590
NONCONST	3998
Street Condition	3473
Illegal Parking	3343
Noise	3321
Traffic Signal Condition	3145
Dirty Conditions	2653
Water System	2636
Noise - Commercial	2578
ELECTRIC	2350
Broken Muni Meter	2070
Noise - Street/Sidewalk	1928
Sanitation Condition	1824
Rodent	1632
Sewer	1627
Taxi Complaint	1227
Consumer Complaint	1227
Damaged Tree	1180
Overgrown Tree/Branches	1083
Graffiti	973
Missed Collection (All Materials)	973
Building/Use	942
Root/Sewer/Sidewalk Condition	836
Derelict Vehicle	803
...	
Posting Advertisement	5
Miscellaneous Categories	5
Fire Alarm - Modification	5
Poison Ivy	5
Internal Code	5
Transportation Provider Complaint	4
Special Natural Area District (SNAD)	4
Ferry Complaint	4
Illegal Animal Sold	4
Fire Alarm - Replacement	3
Invitation	3
Illegal Fireworks	3
Adopt-A-Basket	3
Window Guard	2
Legal Services Provider Complaint	2
Opinion for the Mayor	2
Public Assembly	2
Misc. Comments	2
DFTA Literature Request	2
Highway Sign - Damaged	1
X-Ray Machine/Equipment	1
Municipal Parking Facility	1
DHS Income Savings Requirement	1
Tunnel Condition	1
Open Flame Permit	1
Snow	1
Trans Fat	1
Ferry Permit	1
nwn	1

Stalled Sites

1

Name: Complaint Type, dtype: int64

Now, the top 10 most common complaints:

In [24]:

```
complaint_counts = complaints['Complaint Type'].value_counts()
complaint_counts[:10]
```

Out[24]:


```
HEATING                14200
GENERAL CONSTRUCTION   7471
Street Light Condition  7117
DOF Literature Request  5797
PLUMBING               5373
PAINT - PLASTER        5149
Blocked Driveway       4590
NONCONST               3998
Street Condition       3473
Illegal Parking        3343
Name: Complaint Type, dtype: int64
```

In [25]:

```
complaint_counts[:10].plot(kind='bar')
```

Out[25]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19bf94af98>



Selecting only noise complaints

In [26]:

```
noise_complaints = complaints[complaints['Complaint Type'] == "Noise - Street/Sidewalk"]
noise_complaints[:3]
```

Out[26]:

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	Descriptor
0	26589651	10/31/2013 02:08:41 AM	NaN	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Talking
16	26594086	10/31/2013 12:54:03 AM	10/31/2013 02:16:39 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Music/Party
25	26591573	10/31/2013 12:35:18 AM	10/31/2013 02:41:35 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Talking

3 rows × 52 columns



Why does it work?

In [27]:

```
complaints['Complaint Type'] == "Noise - Street/Sidewalk"
```

Out[27]:

0	True
1	False
2	False
3	False
4	False
5	False
6	False
7	False
8	False
9	False
10	False
11	False
12	False
13	False
14	False
15	False
16	True
17	False
18	False
19	False
20	False
21	False
22	False
23	False
24	False
25	True
26	False
27	False
28	True
29	False
	...
111039	False
111040	False
111041	False
111042	True
111043	False
111044	True
111045	False
111046	False
111047	False
111048	True
111049	False
111050	False
111051	False
111052	False
111053	False
111054	True
111055	False
111056	False
111057	False
111058	False
111059	True
111060	False
111061	False
111062	False
111063	False
111064	False
111065	False
111066	True
111067	False

111067 False

111068 False

Name: Complaint Type, dtype: bool

- when we index our dataframe with this array, we get just the rows where our boolean array evaluated to True
- important - for row filtering by a boolean array the length of our dataframe's index must be the same length as the boolean array used for filtering

More than one condition may be combined with the & operator:


In [28]:

```
is_noise = complaints['Complaint Type'] == "Noise - Street/Sidewalk"
in_brooklyn = complaints['Borough'] == "BROOKLYN"
complaints[is_noise & in_brooklyn][:5]
```

Out[28]:

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	Description
31	26595564	10/31/2013 12:30:36 AM	NaN	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Music/Pa
49	26595553	10/31/2013 12:05:10 AM	10/31/2013 02:43:43 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Talking
109	26594653	10/30/2013 11:26:32 PM	10/31/2013 12:18:54 AM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Music/Pa
236	26591992	10/30/2013 10:02:58 PM	10/30/2013 10:23:20 PM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Talking
370	26594167	10/30/2013 08:38:25 PM	10/30/2013 10:26:28 PM	NYPD	New York City Police Department	Noise - Street/Sidewalk	Loud Music/Pa

5 rows × 52 columns



In [29]:

```
complaints[is_noise & in_brooklyn][['Complaint Type', 'Borough', 'Created Date',
'Descriptor']][:10]
```

Out[29]:

	Complaint Type	Borough	Created Date	Descriptor
31	Noise - Street/Sidewalk	BROOKLYN	10/31/2013 12:30:36 AM	Loud Music/Party
49	Noise - Street/Sidewalk	BROOKLYN	10/31/2013 12:05:10 AM	Loud Talking
109	Noise - Street/Sidewalk	BROOKLYN	10/30/2013 11:26:32 PM	Loud Music/Party
236	Noise - Street/Sidewalk	BROOKLYN	10/30/2013 10:02:58 PM	Loud Talking
370	Noise - Street/Sidewalk	BROOKLYN	10/30/2013 08:38:25 PM	Loud Music/Party
378	Noise - Street/Sidewalk	BROOKLYN	10/30/2013 08:32:13 PM	Loud Talking
656	Noise - Street/Sidewalk	BROOKLYN	10/30/2013 06:07:39 PM	Loud Music/Party
1251	Noise - Street/Sidewalk	BROOKLYN	10/30/2013 03:04:51 PM	Loud Talking
5416	Noise - Street/Sidewalk	BROOKLYN	10/29/2013 10:07:02 PM	Loud Talking
5584	Noise - Street/Sidewalk	BROOKLYN	10/29/2013 08:15:59 PM	Loud Music/Party

Borough with the most noise complaints

In [30]:

```
is_noise = complaints['Complaint Type'] == "Noise - Street/Sidewalk"
noise_complaints = complaints[is_noise]
noise_complaints['Borough'].value_counts()
```

Out[30]:

```
MANHATTAN      917
BROOKLYN       456
BRONX          292
QUEENS         226
STATEN ISLAND   36
Unspecified      1
Name: Borough, dtype: int64
```

We want to normalize the above results by the total number of complaints in a given borough:

In [31]:

```
noise_complaint_counts = noise_complaints['Borough'].value_counts()  
complaint_counts = complaints['Borough'].value_counts()  
noise_complaint_counts / complaint_counts
```

Out[31]:


```
BRONX          0.014833  
BROOKLYN       0.013864  
MANHATTAN      0.037755  
QUEENS         0.010143  
STATEN ISLAND  0.007474  
Unspecified    0.000141  
Name: Borough, dtype: float64
```

In [32]:

```
(noise_complaint_counts / complaint_counts).plot(kind='bar')
```

Out[32]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19bf94c8d0>



So Manhattan really does complain more about noise than the other boroughs.

Cleaning messy data

In [33]:

```
complaints['Incident Zip'].unique()
```

Out[33]:

```

array(['11432', '11378', '10032', '10023', '10027', '11372', '1141
9',
      '11417', '10011', '11225', '11218', '10003', '10029', '1046
6',
      '11219', '10025', '10310', '11236', nan, '10033', '11216', '1
0016',
      '10305', '10312', '10026', '10309', '10036', '11433', '1123
5',
      '11213', '11379', '11101', '10014', '11231', '11234', '1045
7',
      '10459', '10465', '11207', '10002', '10034', '11233', '1045
3',
      '10456', '10469', '11374', '11221', '11421', '11215', '1000
7',
      '10019', '11205', '11418', '11369', '11249', '10005', '1000
9',
      '11211', '11412', '10458', '11229', '10065', '10030', '1122
2',
      '10024', '10013', '11420', '11365', '10012', '11214', '1121
2',
      '10022', '11232', '11040', '11226', '10281', '11102', '1120
8',
      '10001', '10472', '11414', '11223', '10040', '11220', '1137
3',
      '11203', '11691', '11356', '10017', '10452', '10280', '1121
7',
      '10031', '11201', '11358', '10128', '11423', '10039', '1001
0',
      '11209', '10021', '10037', '11413', '11375', '11238', '1047
3',
      '11103', '11354', '11361', '11106', '11385', '10463', '1046
7',
      '11204', '11237', '11377', '11364', '11434', '11435', '1121
0',
      '11228', '11368', '11694', '10464', '11415', '10314', '1030
1',
      '10018', '10038', '11105', '11230', '10468', '11104', '1047
1',
      '11416', '10075', '11422', '11355', '10028', '10462', '1030
6',
      '10461', '11224', '11429', '10035', '11366', '11362', '1120
6',
      '10460', '10304', '11360', '11411', '10455', '10475', '1006
9',
      '10303', '10308', '10302', '11357', '10470', '11367', '1137
0',
      '10454', '10451', '11436', '11426', '10153', '11004', '1142
8',
      '11427', '11001', '11363', '10004', '10474', '11430', '1000
0',
      '10307', '11239', '10119', '10006', '10048', '11697', '1169
2',
      '11693', '10573', '00083', '11559', '10020', '77056', '1177
6',
      '70711', '10282', '11109', '10044', '02061', '77092-2016', '1
4225',
      '55164-0737', '19711', '07306', '000000', 'NO CLUE', '90010',
      '11747', '23541', '11788', '07604', '10112', '11563', '1158
0',
      '07087', '11012', '07003', '11501', '02123', '000000', '1157

```

```

5',
'07109', '11797', '10803', '11716', '11722', '11549-3650', '1
0162',
'23502', '11518', '07020', '08807', '11577', '07114', '1100
3',
'07201', '61702', '10103', '29616-0759', '35209-3114', '1152
0',
'11735', '10129', '11005', '41042', '11590', '06901', '0720
8',
'11530', '13221', '10954', '11111', '10107'], dtype=object)

```

Some of the problems:

- 'NO CLUE', 'N/A', 29616-0759, 83 and NaN values
- some codes have been parsed as strings, some as floats

What we can do:

- normalize 'NO CLUE' and 'N/A' into regular NaN values
- make everything strings
- look at the 83 and other strange codes, and decide what to do

Fixing NaNs and string/float confusion

In [34]:

```

na_values = ['NO CLUE', 'N/A', '0']
complaints = pd.read_csv('data/311-service-requests.csv', na_values=na_values, d
type={'Incident Zip': str})

```

In [35]:

```
complaints['Incident Zip'].unique()
```

array(['11432', '11378', '10032', '10023', '10027', '11372', '11419', '11417', '10011', '11225', '11218', '10003', '10029', '10466', '11219', '10025', '10310', '11236', nan, '10033', '11216', '10016', '10305', '10312', '10026', '10309', '10036', '11433', '11235', '11213', '11379', '11101', '10014', '11231', '11234', '10457', '10459', '10465', '11207', '10002', '10034', '11233', '10453', '10456', '10469', '11374', '11221', '11421', '11215', '10007', '10019', '11205', '11418', '11369', '11249', '10005', '10009', '11211', '11412', '10458', '11229', '10065', '10030', '11222', '10024', '10013', '11420', '11365', '10012', '11214', '11212', '10022', '11232', '11040', '11226', '10281', '11102', '11208', '10001', '10472', '11414', '11223', '10040', '11220', '11373', '11203', '11691', '11356', '10017', '10452', '10280', '11217', '10031', '11201', '11358', '10128', '11423', '10039', '10010', '11209', '10021', '10037', '11413', '11375', '11238', '10473', '11103', '11354', '11361', '11106', '11385', '10463', '10467', '11204', '11237', '11377', '11364', '11434', '11435', '11210', '11228', '11368', '11694', '10464', '11415', '10314', '10301', '10018', '10038', '11105', '11230', '10468', '11104', '10471', '11416', '10075', '11422', '11355', '10028', '10462', '10306', '10461', '11224', '11429', '10035', '11366', '11362', '11206', '10460', '10304', '11360', '11411', '10455', '10475', '10069', '10303', '10308', '10302', '11357', '10470', '11367', '11370', '10454', '10451', '11436', '11426', '10153', '11004', '11428', '11427', '11001', '11363', '10004', '10474', '11430', '10000', '10307', '11239', '10119', '10006', '10048', '11697', '11692', '11693', '10573', '00083', '11559', '10020', '77056', '11776', '70711', '10282', '11109', '10044', '02061', '77092-2016', '14225', '55164-0737', '19711', '07306', '000000', '90010', '11747', '23541', '11788', '07604', '10112', '11563', '11580', '07087', '11042']
--

```

^ ,
    '07093', '11501', '92123', '00000', '11575', '07109', '1179
7',
    '10803', '11716', '11722', '11549-3650', '10162', '23502', '1
1518',
    '07020', '08807', '11577', '07114', '11003', '07201', '6170
2',
    '10103', '29616-0759', '35209-3114', '11520', '11735', '1012
9',
    '11005', '41042', '11590', '06901', '07208', '11530', '1322
1',
    '10954', '11111', '10107'], dtype=object)

```

Zip codes with the dashes

In [36]:

```

rows_with_dashes = complaints['Incident Zip'].str.contains('-').fillna(False)
len(complaints[rows_with_dashes])

```

Out[36]:

5

In [37]:

```
complaints[rows_with_dashes]
```

Out[37]:

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	Descriptor
29136	26550551	10/24/2013 06:16:34 PM	NaN	DCA	Department of Consumer Affairs	Consumer Complaint	False Advertising
30939	26548831	10/24/2013 09:35:10 AM	NaN	DCA	Department of Consumer Affairs	Consumer Complaint	Harassmen
70539	26488417	10/15/2013 03:40:33 PM	NaN	TLC	Taxi and Limousine Commission	Taxi Complaint	Driver Complaint
85821	26468296	10/10/2013 12:36:43 PM	10/26/2013 01:07:07 AM	DCA	Department of Consumer Affairs	Consumer Complaint	Debt Not Owed
89304	26461137	10/09/2013 05:23:46 PM	10/25/2013 01:06:41 AM	DCA	Department of Consumer Affairs	Consumer Complaint	Harassmen

5 rows × 52 columns

In [38]:

```
long_zip_codes = complaints['Incident Zip'].str.len() > 5
complaints['Incident Zip'][long_zip_codes].unique()
```

Out[38]:

```
array(['77092-2016', '55164-0737', '000000', '11549-3650', '29616-0759',
      '35209-3114'], dtype=object)
```

As far as the long zip codes are concerned, they are normal for the USA, but may be truncated to 5 digits:

In [39]:

```
complaints['Incident Zip'] = complaints['Incident Zip'].str.slice(0, 5)
```


83 zip code

Central Park's zip code is 00083 - this code is correct!

00000 zip code


In [40]:

```
complaints[complaints['Incident Zip'] == '00000']
```

Out[40]:

	Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	Descriptor	Location
42600	26529313	10/22/2013 02:51:06 PM	NaN	TLC	Taxi and Limousine Commission	Taxi Complaint	Driver Complaint	Na
60843	26507389	10/17/2013 05:48:44 PM	NaN	TLC	Taxi and Limousine Commission	Taxi Complaint	Driver Complaint	St

2 rows × 9 columns



This entries look bad. So we replace 00000 by NaNs:

In [41]:

```
zero_zips = complaints['Incident Zip'] == '00000'
complaints.loc[zero_zips, 'Incident Zip'] = np.nan
```

In [42]:

```
unique_zips = complaints['Incident Zip'].sort_values().unique()
```

In [43]:

```
unique_zips
```

Out[43]:

```

array(['00083', '02061', '06901', '07020', '07087', '07093', '0710
9',
      '07114', '07201', '07208', '07306', '07604', '08807', '1000
0',
      '10001', '10002', '10003', '10004', '10005', '10006', '1000
7',
      '10009', '10010', '10011', '10012', '10013', '10014', '1001
6',
      '10017', '10018', '10019', '10020', '10021', '10022', '1002
3',
      '10024', '10025', '10026', '10027', '10028', '10029', '1003
0',
      '10031', '10032', '10033', '10034', '10035', '10036', '1003
7',
      '10038', '10039', '10040', '10044', '10048', '10065', '1006
9',
      '10075', '10103', '10107', '10112', '10119', '10128', '1012
9',
      '10153', '10162', '10280', '10281', '10282', '10301', '1030
2',
      '10303', '10304', '10305', '10306', '10307', '10308', '1030
9',
      '10310', '10312', '10314', '10451', '10452', '10453', '1045
4',
      '10455', '10456', '10457', '10458', '10459', '10460', '1046
1',
      '10462', '10463', '10464', '10465', '10466', '10467', '1046
8',
      '10469', '10470', '10471', '10472', '10473', '10474', '1047
5',
      '10573', '10803', '10954', '11001', '11003', '11004', '1100
5',
      '11040', '11042', '11101', '11102', '11103', '11104', '1110
5',
      '11106', '11109', '11111', '11201', '11203', '11204', '1120
5',
      '11206', '11207', '11208', '11209', '11210', '11211', '1121
2',
      '11213', '11214', '11215', '11216', '11217', '11218', '1121
9',
      '11220', '11221', '11222', '11223', '11224', '11225', '1122
6',
      '11228', '11229', '11230', '11231', '11232', '11233', '1123
4',
      '11235', '11236', '11237', '11238', '11239', '11249', '1135
4',
      '11355', '11356', '11357', '11358', '11360', '11361', '1136
2',
      '11363', '11364', '11365', '11366', '11367', '11368', '1136
9',
      '11370', '11372', '11373', '11374', '11375', '11377', '1137
8',
      '11379', '11385', '11411', '11412', '11413', '11414', '1141
5',
      '11416', '11417', '11418', '11419', '11420', '11421', '1142
2',
      '11423', '11426', '11427', '11428', '11429', '11430', '1143
2',
      '11433', '11434', '11435', '11436', '11501', '11518', '1152
0']

```

```

    '11530', '11549', '11559', '11563', '11575', '11577', '1158
0',
    '11590', '11691', '11692', '11693', '11694', '11697', '1171
6',
    '11722', '11735', '11747', '11776', '11788', '11797', '1322
1',
    '14225', '19711', '23502', '23541', '29616', '35209', '4104
2',
    '55164', '61702', '70711', '77056', '77092', '90010', '9212
3', nan], dtype=object)

```

Are all complaints from NY?

For the sake of simplicity we can assume that the zips starting with '0' and '1' are okay (there are some exceptions though):

In [44]:

```

zips = complaints['Incident Zip']
is_close = zips.str.startswith('0') | zips.str.startswith('1')
is_far = ~(is_close) & zips.notnull()

```

In [45]:

```
zips[is_far]
```

Out[45]:

```

12102    77056
13450    70711
29136    77092
30939    55164
44008    90010
47048    23541
57636    92123
71001    92123
71834    23502
80573    61702
85821    29616
89304    35209
94201    41042

```

Name: Incident Zip, dtype: object

In [46]:

```
complaints[is_far][['Incident Zip', 'Descriptor', 'City']].sort_values('Incident Zip')
```

Out[46]:

	Incident Zip	Descriptor	City
71834	23502	Harassment	NORFOLK
47048	23541	Harassment	NORFOLK
85821	29616	Debt Not Owed	GREENVILLE
89304	35209	Harassment	BIRMINGHAM
94201	41042	Harassment	FLORENCE
30939	55164	Harassment	ST. PAUL
80573	61702	Billing Dispute	BLOOMINGTON
13450	70711	Contract Dispute	CLIFTON
12102	77056	Debt Not Owed	HOUSTON
29136	77092	False Advertising	HOUSTON
44008	90010	Billing Dispute	LOS ANGELES
57636	92123	Harassment	SAN DIEGO
71001	92123	Billing Dispute	SAN DIEGO

There are really requests coming from LA and Houston!

String operations

In [47]:

```
weather_2012 = pd.read_csv('data/weather_2012.csv', parse_dates=True,
index_col='Date/Time')
weather_2012[:5]
```

Out[47]:

	Temp (C)	Dew Point Temp (C)	Rel Hum (%)	Wind Spd (km/h)	Visibility (km)	Stn Press (kPa)	Weather
Date/Time							
2012-01-01 00:00:00	-1.8	-3.9	86	4	8.0	101.24	Fog
2012-01-01 01:00:00	-1.8	-3.7	87	4	8.0	101.24	Fog
2012-01-01 02:00:00	-1.8	-3.4	89	7	4.0	101.26	Freezing Drizzle,Fog
2012-01-01 03:00:00	-1.5	-3.2	88	6	4.0	101.27	Freezing Drizzle,Fog
2012-01-01 04:00:00	-1.5	-3.3	88	7	4.8	101.23	Fog

Looking for text

In [48]:

```
weather_description = weather_2012['Weather']
is_snowing = weather_description.str.contains('Snow')
is_snowing[:5] #this is not very useful
```

Out[48]:

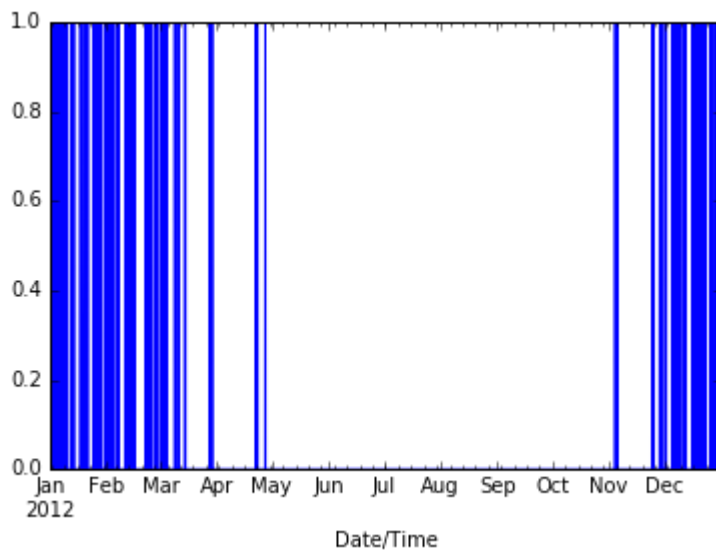
```
Date/Time
2012-01-01 00:00:00    False
2012-01-01 01:00:00    False
2012-01-01 02:00:00    False
2012-01-01 03:00:00    False
2012-01-01 04:00:00    False
Name: Weather, dtype: bool
```

In [49]:

```
is_snowing.plot() #more useful
```

Out[49]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19bc501eb8>




Median temperature in each month

In [50]:

```
weather_2012['Temp (C)'].resample('M').apply(np.median).plot(kind='bar')
```

Out[50]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19bc2edb70>



The snowiest month

We can think of snowiness as being a bunch of 1s and 0s instead of Trues and Falses:

In [51]:

```
is_snowing.astype(float)[:10]
```

Out[51]:

```
Date/Time
2012-01-01 00:00:00    0.0
2012-01-01 01:00:00    0.0
2012-01-01 02:00:00    0.0
2012-01-01 03:00:00    0.0
2012-01-01 04:00:00    0.0
2012-01-01 05:00:00    0.0
2012-01-01 06:00:00    0.0
2012-01-01 07:00:00    0.0
2012-01-01 08:00:00    0.0
2012-01-01 09:00:00    0.0
Name: Weather, dtype: float64
```

Then we use resample to find the percentage of time it was snowing each month:

In [52]:

```
is_snowing.astype(float).resample('M').apply(np.mean)
```

Out[52]:


```
Date/Time
2012-01-31    0.240591
2012-02-29    0.162356
2012-03-31    0.087366
2012-04-30    0.015278
2012-05-31    0.000000
2012-06-30    0.000000
2012-07-31    0.000000
2012-08-31    0.000000
2012-09-30    0.000000
2012-10-31    0.000000
2012-11-30    0.038889
2012-12-31    0.251344
Freq: M, Name: Weather, dtype: float64
```


In [53]:

```
is_snowing.astype(float).resample('M').apply(np.mean).plot(kind='bar')
```

Out[53]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19bc191da0>



- in 2012, December was the snowiest month
- it starts snowing pretty abruptly in November, and then tapers off slowly and takes a long time to stop, with the last snow usually being in April or May

Temperature and snowiness - summary

In [54]:

```
temperature = weather_2012['Temp (C)'].resample('M').apply(np.median)
is_snowing = weather_2012['Weather'].str.contains('Snow')
snowiness = is_snowing.astype(float).resample('M').apply(np.mean)

temperature.name = "Temperature"
snowiness.name = "Snowiness"
```

We can use concat to combine the two statistics into a single dataframe:

In [55]:

```
stats = pd.concat([temperature, snowiness], axis=1)
stats
```

Out[55]:


	Temperature	Snowiness
Date/Time		
2012-01-31	-7.05	0.240591
2012-02-29	-4.10	0.162356
2012-03-31	2.60	0.087366
2012-04-30	6.30	0.015278
2012-05-31	16.05	0.000000
2012-06-30	19.60	0.000000
2012-07-31	22.90	0.000000
2012-08-31	22.20	0.000000
2012-09-30	16.10	0.000000
2012-10-31	11.30	0.000000
2012-11-30	1.05	0.038889
2012-12-31	-2.85	0.251344

In [56]:

```
stats.plot(kind='bar') #the scales are wrong
```

Out[56]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f19bc12def0>




In [57]:

```
stats.plot(kind='bar', subplots=True, figsize=(15, 10))
```

Out[57]:

```
array([<matplotlib.axes._subplots.AxesSubplot object at 0x7f19bc241b70>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f19b66d6438>], dtype=object)
```



Parsing Unix timestamps

In [58]:

```
popcon = pd.read_csv('data/popularity-contest.txt', sep=' ', )[:-1]
popcon.columns = ['atime', 'ctime', 'package-name', 'mru-program', 'tag']
popcon[:5]
```

Out[58]:

	atime	ctime	package-name	mru-program	tag
0	1387295797	1367633260	perl-base	/usr/bin/perl	NaN
1	1387295796	1354370480	login	/bin/su	NaN
2	1387295743	1354341275	libtalloc2	/usr/lib/x86_64-linux-gnu/libtalloc.so.2.0.7	NaN
3	1387295743	1387224204	libwbclient0	/usr/lib/x86_64-linux-gnu/libwbclient.so.0	<RECENT-CTIME>
4	1387295742	1354341253	libselsinux1	/lib/x86_64-linux-gnu/libselsinux.so.1	NaN

In [59]:

```
popcon['atime'] = popcon['atime'].astype(int)
popcon['ctime'] = popcon['ctime'].astype(int)
```

- internally, numpy datetimes are already stored as Unix timestamps!!!
- all we need to do is to tell pandas that these integers are actually datetimes → it doesn't need to do any conversion at all

In [60]:

```
popcon['atime'] = pd.to_datetime(popcon['atime'], unit='s')
popcon['ctime'] = pd.to_datetime(popcon['ctime'], unit='s')
```

In [61]:

```
popcon[:5]
```

Out[61]:

	atime	ctime	package-name	mru-program	tag
0	2013-12-17 15:56:37	2013-05-04 02:07:40	perl-base	/usr/bin/perl	NaN
1	2013-12-17 15:56:36	2012-12-01 14:01:20	login	/bin/su	NaN
2	2013-12-17 15:55:43	2012-12-01 05:54:35	libtalloc2	/usr/lib/x86_64-linux-gnu/libtalloc.so.2.0.7	NaN
3	2013-12-17 15:55:43	2013-12-16 20:03:24	libwbclient0	/usr/lib/x86_64-linux-gnu/libwbclient.so.0	<RECENT-CTIME>
4	2013-12-17 15:55:42	2012-12-01 05:54:13	libselenium1	/lib/x86_64-linux-gnu/libselinux.so.1	NaN

Suppose we want to look at all packages that aren't libraries.

First, we want to get rid of everything with timestamp 0:

In [62]:

```
popcon = popcon[popcon['atime'] > '1970-01-01']
```

Now, we look at rows not containing the 'lib' string:

In [63]:

```
nonlibraries = popcon[~popcon['package-name'].str.contains('lib')]
nonlibraries.sort_values(by='ctime', ascending=False)[:10]
```

Out[63]:

	atime	ctime	package-name	mru-program	tag
57	2013-12-17 04:55:39	2013-12-17 04:55:42	ddd	/usr/bin/ddd	<RECECTIME:
450	2013-12-16 20:03:20	2013-12-16 20:05:13	nodejs	/usr/bin/npm	<RECECTIME:
454	2013-12-16 20:03:20	2013-12-16 20:05:04	switchboard-plug-keyboard	/usr/lib/plugs/pantheon/keyboard/options.txt	<RECECTIME:
445	2013-12-16 20:03:20	2013-12-16 20:05:04	thunderbird-locale-en	/usr/lib/thunderbird-addons/extensions/langpac...	<RECECTIME:
396	2013-12-16 20:08:27	2013-12-16 20:05:03	software-center	/usr/sbin/update-software-center	<RECECTIME:
449	2013-12-16 20:03:20	2013-12-16 20:05:00	samba-common-bin	/usr/bin/net.samba3	<RECECTIME:
397	2013-12-16 20:08:25	2013-12-16 20:04:59	postgresql-client-9.1	/usr/lib/postgresql/9.1/bin/psql	<RECECTIME:
398	2013-12-16 20:08:23	2013-12-16 20:04:58	postgresql-9.1	/usr/lib/postgresql/9.1/bin/postmaster	<RECECTIME:
452	2013-12-16 20:03:20	2013-12-16 20:04:55	php5-dev	/usr/include/php5/main/snprintf.h	<RECECTIME:
440	2013-12-16 20:03:20	2013-12-16 20:04:54	php-pear	/usr/share/php/XML/Util.php	<RECECTIME:

Loading data from SQL

In [64]:

```
import sqlite3
```

In [65]:

```
con = sqlite3.connect("data/weather_2012.sqlite")
df = pd.read_sql("SELECT * from weather_2012 LIMIT 3", con)
df
```

Out[65]:

	id	date_time	temp
0	1	2012-01-01 00:00:00	-1.8
1	2	2012-01-01 01:00:00	-1.8
2	3	2012-01-01 02:00:00	-1.8

In [66]:

```
df = pd.read_sql("SELECT * from weather_2012 LIMIT 3", con, index_col='id') #indexed by 'id' column
df
```

Out[66]:

	date_time	temp
id		
1	2012-01-01 00:00:00	-1.8
2	2012-01-01 01:00:00	-1.8
3	2012-01-01 02:00:00	-1.8

In [67]:

```
df = pd.read_sql("SELECT * from weather_2012 LIMIT 3", con, index_col=['id', 'date_time']) #indexed by multiple columns
df
```

Out[67]:

		temp
id	date_time	
1	2012-01-01 00:00:00	-1.8
2	2012-01-01 01:00:00	-1.8
3	2012-01-01 02:00:00	-1.8

Pandas vs SQL

In [68]:

```
url = 'https://raw.githubusercontent.com/pydata/pandas/master/pandas/tests/data/tips.csv'
tips = pd.read_csv(url)
tips.head()
```

Out[68]:

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

SELECT

```
SELECT total_bill, tip, smoker, time
FROM tips
LIMIT 5;
```

In [69]:

```
tips[['total_bill', 'tip', 'smoker', 'time']].head(5)
```

Out[69]:

	total_bill	tip	smoker	time
0	16.99	1.01	No	Dinner
1	10.34	1.66	No	Dinner
2	21.01	3.50	No	Dinner
3	23.68	3.31	No	Dinner
4	24.59	3.61	No	Dinner

WHERE

```
SELECT *
FROM tips
WHERE time = 'Dinner'
LIMIT 5;
```


In [70]:

```
tips[tips['time'] == 'Dinner'].head(5)
```

Out[70]:

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

```
-- tips greater than 5$ at dinner
SELECT *
FROM tips
WHERE time = 'Dinner' AND tip > 5.00;
```

In [71]:

```
tips[(tips['time'] == 'Dinner') & (tips['tip'] > 5.00)]
```

Out[71]:

	total_bill	tip	sex	smoker	day	time	size
23	39.42	7.58	Male	No	Sat	Dinner	4
44	30.40	5.60	Male	No	Sun	Dinner	4
47	32.40	6.00	Male	No	Sun	Dinner	4
52	34.81	5.20	Female	No	Sun	Dinner	4
59	48.27	6.73	Male	No	Sat	Dinner	4
116	29.93	5.07	Male	No	Sun	Dinner	4
155	29.85	5.14	Female	No	Sun	Dinner	5
170	50.81	10.00	Male	Yes	Sat	Dinner	3
172	7.25	5.15	Male	Yes	Sun	Dinner	2
181	23.33	5.65	Male	Yes	Sun	Dinner	2
183	23.17	6.50	Male	Yes	Sun	Dinner	4
211	25.89	5.16	Male	Yes	Sat	Dinner	4
212	48.33	9.00	Male	No	Sat	Dinner	4
214	28.17	6.50	Female	Yes	Sat	Dinner	3
239	29.03	5.92	Male	No	Sat	Dinner	3

```
-- tip at dinner, group size >= 5, total bill > 45
SELECT *
FROM tips
WHERE size >= 5 OR total_bill > 45;
```

In [72]:

```
tips[(tips['size'] >= 5) | (tips['total_bill'] > 45)]
```

Out[72]:

	total_bill	tip	sex	smoker	day	time	size
59	48.27	6.73	Male	No	Sat	Dinner	4
125	29.80	4.20	Female	No	Thur	Lunch	6
141	34.30	6.70	Male	No	Thur	Lunch	6
142	41.19	5.00	Male	No	Thur	Lunch	5
143	27.05	5.00	Female	No	Thur	Lunch	6
155	29.85	5.14	Female	No	Sun	Dinner	5
156	48.17	5.00	Male	No	Sun	Dinner	6
170	50.81	10.00	Male	Yes	Sat	Dinner	3
182	45.35	3.50	Male	Yes	Sun	Dinner	3
185	20.69	5.00	Male	No	Sun	Dinner	5
187	30.46	2.00	Male	Yes	Sun	Dinner	5
212	48.33	9.00	Male	No	Sat	Dinner	4
216	28.15	3.00	Male	Yes	Sat	Dinner	5

GROUPBY

```
SELECT sex, count(*)
FROM tips
GROUP BY sex;
```

In [73]:

```
tips.groupby('sex').size()
```

Out[73]:

```
sex
Female      87
Male       157
dtype: int64
```

In [74]:

```
tips.groupby('sex').count() #non-zero entries in each column
```

Out[74]:

	total_bill	tip	smoker	day	time	size
sex						
Female	87	87	87	87	87	87
Male	157	157	157	157	157	157

In [75]:

```
tips.groupby('sex')['total_bill'].count()
```

Out[75]:

```
sex
Female      87
Male       157
Name: total_bill, dtype: int64
```

```
SELECT day, AVG(tip), COUNT(*)
FROM tips
GROUP BY day;
```

In [76]:

```
tips.groupby('day').agg({'tip': np.mean, 'day': np.size})
```

Out[76]:

	day	tip
day		
Fri	19	2.734737
Sat	87	2.993103
Sun	76	3.255132
Thur	62	2.771452

```
SELECT smoker, day, COUNT(*), AVG(tip)
FROM tips
GROUP BY smoker, day;
```